

```
In [1]: import os
os.chdir(r"C:\Users\JAHNAVI\Documents\finance_categorization_project")
print("Working directory:", os.getcwd())
Working directory: C:\Users\JAHNAVI\Documents\finance_categorization_project
```

```
In [2]: import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import sqlalchemy as sa
from sqlalchemy import create_engine
from sqlalchemy.engine import URL
import pyodbc
import pypyodbc
from rapidfuzz import process
```

```
In [3]: import os
import joblib
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import (
    OneHotEncoder,
    FunctionTransformer,
    MinMaxScaler
)
from sklearn.impute import SimpleImputer

from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import (
    accuracy_score,
    classification_report,
    confusion_matrix
)
```

```
In [4]: #!pip install rapidfuzz
```

DB Connection

```
In [5]: conn = f"""
    DRIVER={"ODBC Driver 17 for SQL Server"};
    SERVER={"DESKTOP-62UF0V4"};
    INSTANCENAME="\JAHNAVI";
    port=1433;
    DATABASE={"Budget_Transactions"};
    Trusted_Connection=yes;
"""
```

```
In [6]: con_str = URL.create('mssql+pyodbc', query={'odbc_connect': conn})
```

```
In [7]: engine=create_engine(con_str, module=pypyodbc, fast_executemany=True)
```

```
In [8]: df = pd.read_sql("SELECT * FROM financial_transactions", engine)
print(df.head())
```

	transaction_id	user_id	date	transaction_type	category	amount	\
0	T4999	U018	2023-04-25	Expense	Educator	3888	
1	T12828	U133	08/05/2022	Expense	rent	649	
2	T7403	U091	31-12-23	Income	Freelance	13239	
3	T12350	U097	None	Expense	Fod	6299	
4	T7495	U088	10/28/2022	Expense	entertainment	2287	
	payment_mode	location	notes				
0	card	Ahmedabad	Movie tickets				
1	None	Hyderabad	asdfgh				
2	Csh	BAN	Books				
3	Bank Transfer	AHMEDABAD	Electricity bill				
4	CARD	Hyderabad	None				

Data Cleaning

```
In [9]: df = df.drop_duplicates()
```

```
In [10]: df.head()
```

```
Out[10]:   transaction_id  user_id      date  transaction_type  category  amount  payment_mode  location  notes
0          T4999    U018  2023-04-25     Expense  Educaton    3888       card  Ahmedabad Movie tickets
1          T12828    U133  08/05/2022     Expense        rent     649      None  Hyderabad  asdfgh
2          T7403    U091  31-12-23      Income  Freelance   13239       Csh      BAN    Books
3          T12350    U097      None     Expense        Fod    6299  Bank Transfer  AHMEDABAD Electricity bill
4          T7495    U088  10/28/2022     Expense entertainment   2287      CARD  Hyderabad      None
```

```
In [11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 30032 entries, 0 to 31735
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   transaction_id    30032 non-null   object 
 1   user_id           30032 non-null   object 
 2   date              29243 non-null   object 
 3   transaction_type  30032 non-null   object 
 4   category          29604 non-null   object 
 5   amount             29591 non-null   object 
 6   payment_mode      28958 non-null   object 
 7   location           28242 non-null   object 
 8   notes              26440 non-null   object 
dtypes: object(9)
memory usage: 2.3+ MB
```

```
In [12]: df.isnull().sum()
```

```
Out[12]: transaction_id      0
user_id            0
date              789
transaction_type  0
category          428
amount            441
payment_mode      1074
location           1790
notes             3592
dtype: int64
```

```
In [13]: df['category'] = df['category'].fillna("unknown")
```

```
In [14]: df["category"].value_counts()
```

```
Out[14]: category
Food            4095
Rent            3451
Travel          1781
Utilities        1523
Entertainment   1313
...
Savigns         1
Entertainmnet   1
Salayr          1
Other IIIncome   1
Educatioon       1
Name: count, Length: 239, dtype: int64
```

```
In [15]: df["category"].unique()
```

```
Out[15]: array(['Educaton', 'rent', 'Freelance', 'Fod', 'entertainment', 'Foods',
   'education', 'Salary', 'Utlties', 'Others', 'Utility', 'Utlities',
   'Rentt', 'Food', 'FOOD', 'Travel', 'food', 'health',
   'Entertainment', 'Travl', 'Investment', 'Foodd', 'HEALTH', 'RENT',
   'Helth', 'Education', 'utilities', 'Rent', 'Rnt', 'savings', 'EDU',
   'Traval', 'unknown', 'Bonus', 'travel', 'Utilities', 'Entertain',
   'Saving', 'Entrtnmnt', 'TRAVEL', 'SAVINGS', 'Misc', 'others',
   'Health', 'Other', 'Savings', 'OTHERS', 'Other Income',
   'Entertainmennt', 'Salaryy', 'Reent', 'ood', 'Traavel', 'Tavel',
   'OOthers', 'Utilites', 'eRnt', 'Trave', 'Savingss', 'Ren',
   'rTavel', 'Foo', 'Otehr Income', 'Otehrs', 'Retn', 'Heealth',
   'Traveel', 'Rennt', 'Ret', 'Food', 'Othhers', 'aSvings',
   'Utilitise', 'Travle', 'Healt', 'FFood', 'Bonu', 'BBonus', 'Bons',
   'oFod', 'Eduation', 'Entertainmtn', 'Otherrs', 'Salry', 'Saary',
   'Utilitiees', 'Entertainemnt', 'thers', 't0thers', 'Utilties',
   'RRent', 'ravel', 'Sallary', 'Savnigs', 'Trravel', 'Utilties',
   'Heaalth', 'Rnet', 'Saviings', 'Bonus', 'ent', 'Saivngs',
   'Savngs', 'Entetainment', 'Svings', 'Salaary', 'Entertaniment',
   'Utilitis', 'Bouns', 'Utiltiies', 'Oters', 'aSlary',
   't0her Income', 'Enntertainment', 'Fodo', 'Edcation', 'Trael',
   'Enttertainment', 'Uttilities', 'Other Icnome', 'Educationn',
   'Uillties', 'Otthers', 'Othres', 'Ohters', 'Educatino',
   'Other Income', 'Entertianment', 'Enteertainment', 'Savvings',
   'Other Incoome', 'Entertaiment', 'Oter Income', 'Utliities',
   'Utiilities', 'Savinngs', 'nErtainment', 'Entertainment',
   'Etertainment', 'Trvel', 'Savigs', 'Utiities', 'Entretainment',
   'Uitilities', 'Bous', 'Enterertainment', 'Entertainment', 'Saavngs',
   'Othrs', 'Othe Income', 'onus', 'Utiliteis', 'Tarvel',
   'Entertainmetn', 'Trvael', 'Utilitiies', 'EEntertainment',
   'Helath', 'Travell', 'Utilitiess', 'Salay', 'Bonnus', 'tilities',
   'alary', 'Utilities', 'Othre Income', 'Savins', 'Educcation',
   'Other Incoem', 'Bnous', 'EEducation', 'tUilities', 'oBnus',
   'Traevl', 'SSavings', 'Utilitie', 'Saalry', 'TTtravel',
   'Enteratinment', 'Travvel', 'Entertainmnet', 'Othr Income',
   'Other Inncome', 'Slary', 'Educatio', 'Etnertainment', 'Othes',
   'Entertaainment', 'Entertainmt', 'Othe rIncome', 'Otheers',
   'Other Incom', 'OtherIncome', 'O0ther Income', 'Other nIcome',
   'Educatiion', 'Entertainment', 'UUtilities', 'Eudcation',
   'Entertaiinment', 'SSalary', 'Bnus', 'Bonsu', 'Boonus',
   'Eduucation', 'Ohers', 'Eucation', 'Enertainment', 'Savigns',
   'avings', 'Saings', 'Salayr', 'Other IIcome', 'Salary',
   'Savinggs', 'Saalary', 'Other Inome', 'Ohter Income',
   'Entetainment', 'Entrtainment', 'Othesn', 'Other Icome',
   'Entertainment', 'Edducation', 'Entertainnmt', 'Uilities',
   'Heath', 'Otheer Income', 'ealth', 'Salray', 'Other Incmoe',
   'Oher Income', 'Entertainmet', 'Healh', 'Other Incoe', 'Savinsg',
   'Otherss', 'Slaary', 'Educatioon'], dtype=object)
```

```
In [16]: df["category"] = df["category"].str.lower().str.strip()
```

```
In [17]: df["category"].value_counts().head(30)
```

```
Out[17]: category
food           5188
rent           4492
travel          2529
utilities        1813
entertainment    1620
others           1176
bonus            1125
salary            1111
savings           843
education         837
other income      707
health             657
fod               555
foods              541
foood              516
rentt              511
rnt                503
freelance          463
investment          455
unknown             428
traval              386
travl              357
entertain            333
utility              320
entrtnmnt            315
utilties              315
utlities              287
educaton              243
edu                 226
helth                150
Name: count, dtype: int64
```

```
In [18]: correct_categories = [
    "food", "rent", "travel", "utilities", "entertainment",
    "education", "others", "bonus", "salary", "savings",
    "health", "other income", "freelance", "investment", "misc"
]
```

```
In [19]: def fuzzy_fix(x):
    if pd.isna(x):
        return x
    match = process.extractOne(x, correct_categories, score_cutoff=80)
    return match[0] if match else x # if similarity <80, keep original

df["category"] = df["category"].apply(fuzzy_fix)
```

```
In [20]: df["category"].value_counts().head(30)
```

```
Out[20]: category
food           6912
rent           5599
travel          3378
utilities        2502
entertainment    2360
education         1328
others           1272
bonus            1172
salary            1160
savings           976
health             820
other income      743
freelance          463
investment          455
unknown             428
utility              320
misc                44
ernt                23
fodo                23
retn                19
ofod                19
rnet                16
Name: count, dtype: int64
```

```
In [21]: fix = {
    "utility": "utilities",
    "ernt": "rent",
    "retn": "rent",
    "rnet": "rent",
    "fodo": "food",
    "ofod": "food"
}

df["category"] = df["category"].replace(fix)
```

```
In [22]: df["category"].value_counts().head(40)
```

```
Out[22]: category
food            6954
rent            5657
travel          3378
utilities        2822
entertainment   2360
education        1328
others           1272
bonus            1172
salary            1160
savings           976
health             820
other income     743
freelance         463
investment        455
unknown           428
misc                44
Name: count, dtype: int64
```

```
In [23]: df["transaction_id"].isnull().sum()
```

```
Out[23]: 0
```

```
In [24]: df["transaction_id"].nunique(), df.shape[0]
```

```
Out[24]: (24896, 30032)
```

```
In [25]: df["transaction_id"].str.len().value_counts()
```

```
Out[25]: transaction_id
6      19327
5      10705
Name: count, dtype: int64
```

```
In [26]: df["transaction_id"].str.contains(r'^[A-Za-z0-9\-\-]+$', regex=True).value_counts()
```

```
Out[26]: transaction_id
True      30032
Name: count, dtype: int64
```

```
In [27]: print(df["transaction_id"].nunique(), df.shape[0])
print(df["transaction_id"].str.len().value_counts().head(10))
```

```
24896 30032
transaction_id
6      19327
5      10705
Name: count, dtype: int64
```

```
In [28]: dupe_ids = df["transaction_id"].value_counts()
dupe_ids[dupe_ids > 1].head(30)
```

```
Out[28]: transaction_id
T12746      5
T13481      5
T13482      5
T11588      4
T13033      4
T8639       4
T10004      4
T10921      4
T13685      4
T13347      4
T13840      4
T4457       4
T13150      4
T12726      4
T12084      4
T11034      4
T10148       4
T10659      4
T1095       4
T12030      4
T4435       4
T6071       4
T12126      4
T12923      4
T8787       4
T11886      4
T10730      3
T0829       3
T6006       3
T12753      3
Name: count, dtype: int64
```

```
In [29]: #Copying original IDs into a new column
df["clean_transaction_id"] = df["transaction_id"].astype(str)

#Function to add suffix only to duplicates

def fix_dupes(series):
    if len(series) == 1:
        return series
    out = []
    for i, val in enumerate(series):
        out.append(val if i == 0 else f"{val}_{i}")
    return pd.Series(out, index=series.index)

#Detecting duplicates to fix only them

dup_mask = df["transaction_id"].duplicated(keep=False)

df.loc[dup_mask, "clean_transaction_id"] = (
    df.loc[dup_mask]
        .groupby("transaction_id")["transaction_id"]
        .transform(fix_dupes)
)

#checking uniqueness
print("Unique clean IDs:", df["clean_transaction_id"].nunique())
print("Total rows:", df.shape[0])
```

Unique clean IDs: 30032
 Total rows: 30032

```
In [30]: df["date"] = pd.to_datetime(df["date"], errors="coerce")
```

```
In [31]: df["date"].isnull().sum()
```

```
Out[31]: 20082
```

```
In [32]: df["amount"] = df["amount"].astype(str).str.replace(r"\d\.\-\]", "", regex=True)
```

```
In [33]: df["amount"] = pd.to_numeric(df["amount"], errors="coerce")
```

```
In [34]: df["amount"].isna().sum()
```

```
Out[34]: 447
```

```
In [35]: df["transaction_type"].value_counts()
```

```
Out[35]: transaction_type
Expense    25517
Income     4515
Name: count, dtype: int64
```

```
In [36]: df["transaction_type"] = df["transaction_type"].astype(str).str.lower().str.strip()
```

```
In [37]: df["payment_mode"].value_counts()
```

```
Out[37]: payment_mode
UPI            3213
Bank Transfer  3122
Card           3067
Cash           3049
Upi           882
...
715"           1
450"           1
936"           1
997"           1
632"           1
Name: count, Length: 1057, dtype: int64
```

```
In [38]: df["payment_mode"].value_counts().head(30)
```

```
Out[38]: payment_mode
UPI            3213
Bank Transfer  3122
Card           3067
Cash           3049
Upi           882
upi           873
UPi           842
CARD           743
Bank Transfr   739
bank transfer  734
card           725
CASH           724
Csh            724
Crd            721
csh            716
CRD            714
cash           695
BankTransfer   691
Bank_Transfer  685
999"          45
396"          16
219"          14
028"          13
019"          13
278"          13
025"          12
100"          12
079"          12
120"          12
731"          12
Name: count, dtype: int64
```

```
In [39]: df["payment_mode"] = (df["payment_mode"].astype(str).str.lower().str.replace("'", "").str.strip())
```

```
# valid payment modes
valid_modes = ["upi", "bank transfer", "card", "cash"]

#fuzzy function
def fuzzy_pay(x):
    match = process.extractOne(x, valid_modes)
    return match[0] if match and match[1] >= 75 else "other"

# apply fuzzy match
df["payment_mode"] = df["payment_mode"].apply(fuzzy_pay)

print(df["payment_mode"].value_counts())
```

```
payment_mode
other        6229
bank transfer 6013
card         6002
cash         5944
upi          5844
Name: count, dtype: int64
```

```
In [41]: df["location"].isna().sum()
```

```
Out[41]: 1790
```

```
In [42]: df["location"] = df["location"].astype(str).str.lower().str.strip()
```

```
In [43]: df["location"] = df["location"].replace(["nan", "none", "null", ""], None)
```

```
In [44]: df["location"] = df["location"].fillna("unknown")
```

```
In [45]: (
    #contains digits?
    df["location"].str.contains(r"\d", regex=True).any(),
    #contains unwanted characters?
    df["location"].str.contains(r"[^a-zA-Z\s]", regex=True).any()
)
```

```
Out[45]: (False, False)
```

```
In [46]: df["location"].value_counts().head(20)
```

```
Out[46]: location
pune          2030
mumbai        1999
chennai        1999
hyderabad      1988
jaipur         1976
kolkata        1958
delhi          1949
ahmedabad      1903
lucknow         1899
unknown         1790
cash            1284
bank transfer   1218
upi             1208
card             1195
bangalore       1050
bengaluru       968
pun              378
hyd              368
kol              360
del              356
Name: count, dtype: int64
```

```
In [47]: #allowed correct locations
valid_locations = [
    "chennai", "kolkata", "delhi", "bengaluru", "hyderabad",
    "mumbai", "pune", "goa", "agra", "mysore", "ahmedabad", "cochin"
]

#fuzzy match function
def fuzzy_loc(x):
    match = process.extractOne(x, valid_locations)
    return match[0] if match and match[1] >= 70 else "unknown"

#apply fuzzy matching
df["location"] = df["location"].apply(fuzzy_loc)

print(df["location"].value_counts())
```

```
location
unknown        12385
chennai         2697
pune            2408
hyderabad       2360
mumbai          2332
kolkata         2318
delhi           2305
ahmedabad       2259
bengaluru        968
Name: count, dtype: int64
```

```
In [48]: df["user_id"].value_counts().head(20)
```

```
Out[48]: user_id
U117    213
U036    212
U131    209
U149    205
U074    200
U080    199
U113    198
U021    197
U086    197
U075    197
U069    195
U119    194
U130    194
U077    194
U115    193
U141    193
U068    192
U091    192
U062    192
U020    191
Name: count, dtype: int64
```

```
In [49]: df["user_id"].isna().sum()
```

```
Out[49]: 0
```

```
In [50]: df['notes'].value_counts(dropna=False).head(20)
```

```
Out[50]: notes
None            3592
Grocery shopping      1027
Gift              618
Uber to office       602
Dinner at resto       585
Doctor visit          584
Course fee           580
Netflix subscription     579
Salary              567
Monthly rent          565
Paid electricity bill 560
Electricity bill        526
Monthly rent payment     508
Movie tickets          502
Gym membership         499
!!!
ATM withdrawal          495
Coffee              492
Internet bill          487
Restaurant dinner        487
Name: count, dtype: int64
```

```
In [51]: df['notes'] = df['notes'].fillna("unknown")
df['notes'] = df['notes'].replace("None", "unknown")
df['notes'] = df['notes'].replace("", "unknown")
```

```
In [52]: df['notes'] = df['notes'].str.lower()
```

```
In [53]: #Remove special characters
df['notes'] = df['notes'].str.replace(r'^[a-zA-Z\s]', '', regex=True)
```

```
In [54]: #Remove extra spaces
df['notes'] = df['notes'].str.strip()
df['notes'] = df['notes'].str.replace(r'\s+', ' ', regex=True)
```

```
In [55]: #Replace empty values after cleaning
df['notes'] = df['notes'].replace("", "unknown")
```

```
In [56]: df['notes'].value_counts(dropna=False).head(20)
```

```
Out[56]: notes
unknown           4599
grocery shopping   1113
gift                704
uber to office      686
course fee          671
dinner at resto     669
netflix subscription 663
doctor visit         663
monthly rent         655
salary               641
paid electricity bill 639
electricity bill      526
monthly rent payment 508
movie tickets         502
gym membership        499
atm withdrawal        495
coffee                492
restaurant dinner      487
internet bill          487
lunch                  485
Name: count, dtype: int64
```

Machine Learning

In [57]: #Train Baseline Multinomial Naive Bayes Model

```
RANDOM_STATE = 42
TARGET = "category"

#Train/test split
df_model = df.copy()
train_df, test_df = train_test_split(
    df_model, test_size=0.2, random_state=RANDOM_STATE, stratify=df_model[TARGET]
)

X_train = train_df.drop(columns=[TARGET])
y_train = train_df[TARGET]
X_test = test_df.drop(columns=[TARGET])
y_test = test_df[TARGET]

#Date features
def extract_date_features(df_in):
    df2 = df_in.copy()
    dates = pd.to_datetime(df2["date"], errors="coerce")
    df2["dt_year"] = dates.dt.year.fillna(0).astype(int)
    df2["dt_month"] = dates.dt.month.fillna(0).astype(int)
    df2["dt_day"] = dates.dt.day.fillna(0).astype(int)
    df2["dt_dow"] = dates.dt.dayofweek.fillna(0).astype(int)
    return df2[["dt_year", "dt_month", "dt_day", "dt_dow"]]

date_transformer = Pipeline([
    ("extract", FunctionTransformer(func=extract_date_features, validate=False)),
    ("impute", SimpleImputer(strategy="constant", fill_value=0)),
    ("scale", MinMaxScaler()) # FIXED (No negatives)
])

#Feature groups
text_col = "notes"
cat_cols = ["payment_mode", "location", "transaction_type"]
num_cols = ["amount"]

text_transformer = Pipeline([
    ("tfidf", TfidfVectorizer(max_features=20000, ngram_range=(1,2), stop_words="english"))
])

cat_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="constant", fill_value="unknown")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

num_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="median")),
    ("scale", MinMaxScaler()) # FIXED (No negatives)
])

preprocessor = ColumnTransformer(
    transformers=[
        ("text", text_transformer, text_col),
        ("cat", cat_transformer, cat_cols),
        ("num", num_transformer, num_cols),
        ("date", date_transformer, ["date"])
    ],
    remainder="drop"
)

#NB model pipeline
nb = Pipeline([
    ("pre", preprocessor),
    ("clf", MultinomialNB())
])

#Train
nb.fit(X_train, y_train)

#Evaluate
y_pred = nb.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred, zero_division=0))

#Save
os.makedirs("output", exist_ok=True)
joblib.dump(nb, "output/nb_pipeline.joblib")
print("Saved: output/nb_pipeline.joblib")
```

Accuracy: 0.2593640752455469

	precision	recall	f1-score	support
bonus	0.24	0.16	0.19	234
education	0.00	0.00	0.00	266
entertainment	0.00	0.00	0.00	472
food	0.27	0.88	0.41	1391
freelance	0.20	0.12	0.15	93
health	0.00	0.00	0.00	164
investment	0.15	0.07	0.09	91
misc	0.00	0.00	0.00	9
other income	0.32	0.56	0.40	149
others	0.20	0.17	0.18	254
rent	0.22	0.09	0.13	1131
salary	0.26	0.20	0.23	232
savings	0.00	0.00	0.00	195
travel	0.08	0.01	0.02	676
unknown	0.00	0.00	0.00	86
utilities	0.00	0.00	0.00	564
accuracy			0.26	6007
macro avg	0.12	0.14	0.11	6007
weighted avg	0.16	0.26	0.16	6007

Saved: output/nb_pipeline.joblib

In [58]: #Model Evaluation & Error Analysis

```
#Predictions
y_pred = nb.predict(X_test)

#Basic metrics
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred, zero_division=0))

#Confusion Matrix
cm = confusion_matrix(y_test, y_pred, labels=nb.classes_)

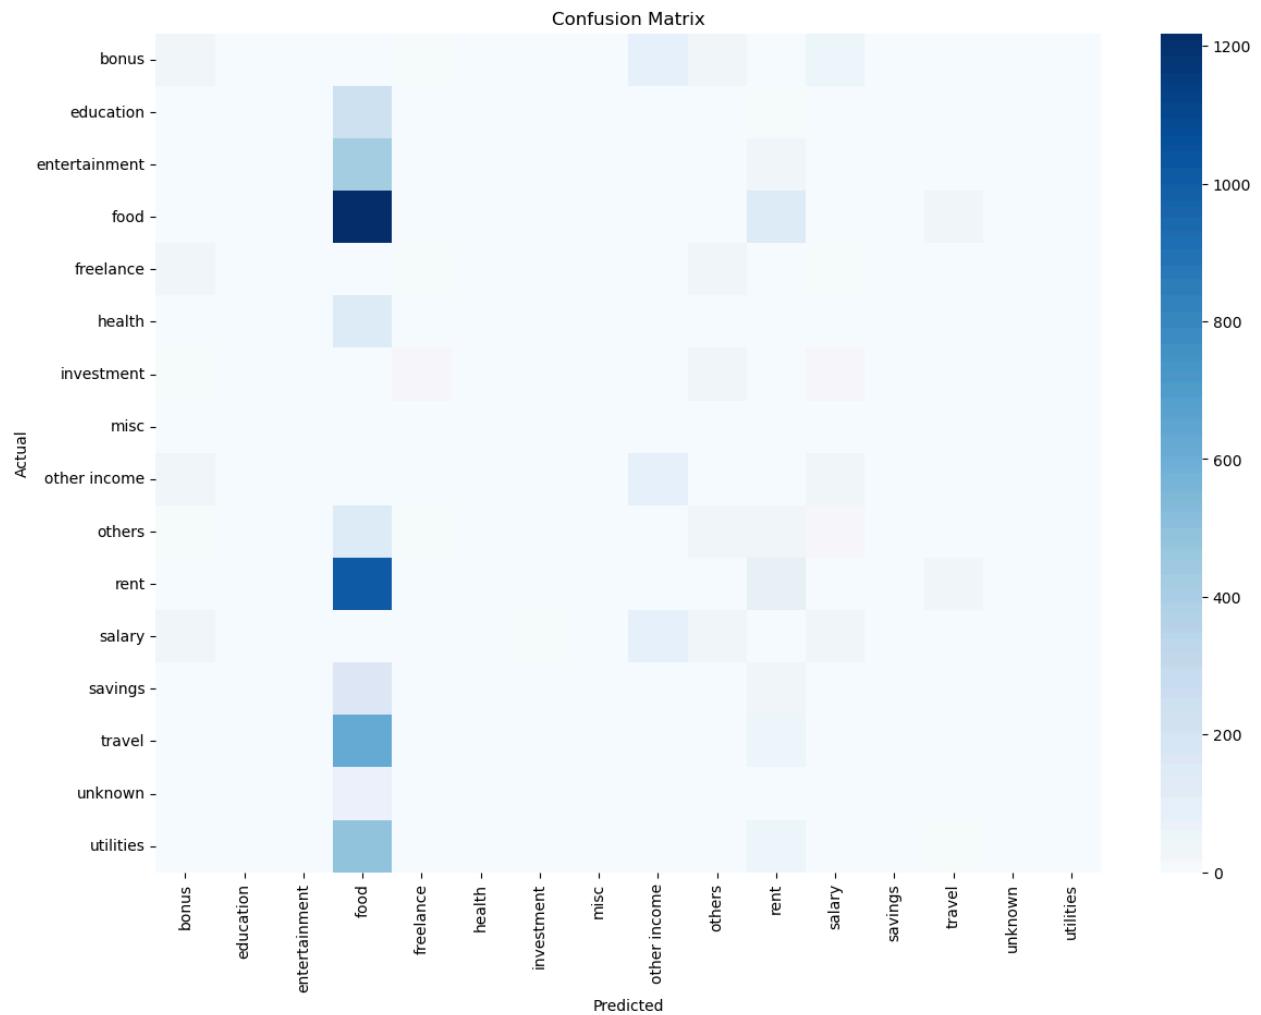
plt.figure(figsize=(14,10))
sns.heatmap(cm, annot=False, cmap="Blues",
            xticklabels=nb.classes_,
            yticklabels=nb.classes_)
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

#Error analysis: show top mistakes
errors = pd.DataFrame({
    "actual": y_test,
    "predicted": y_pred,
    "notes": X_test["notes"]
})

#Only wrong predictions
errors = errors[errors["actual"] != errors["predicted"]]

#Show first 20 mistakes
errors.head(20)
```

	precision	recall	f1-score	support
bonus	0.24	0.16	0.19	234
education	0.00	0.00	0.00	266
entertainment	0.00	0.00	0.00	472
food	0.27	0.88	0.41	1391
freelance	0.20	0.12	0.15	93
health	0.00	0.00	0.00	164
investment	0.15	0.07	0.09	91
misc	0.00	0.00	0.00	9
other income	0.32	0.56	0.40	149
others	0.20	0.17	0.18	254
rent	0.22	0.09	0.13	1131
salary	0.26	0.20	0.23	232
savings	0.00	0.00	0.00	195
travel	0.08	0.01	0.02	676
unknown	0.00	0.00	0.00	86
utilities	0.00	0.00	0.00	564
accuracy			0.26	6007
macro avg	0.12	0.14	0.11	6007
weighted avg	0.16	0.26	0.16	6007



Out[58]:

	actual	predicted	notes
20540	travel	food	monthly rent late
4180	entertainment	food	unknown
8055	entertainment	food	fixed deposit
14714	utilities	food	test
19223	rent	food	grocery shopping via app
9374	travel	food	shopping
8340	rent	food	gym membership
13172	rent	food	uber ride
15750	travel	food	unknown
13335	rent	food	gym membership
17342	bonus	salary	kbtvjj3opj9
1786	rent	food	unknown
22264	food	rent	salary via app
17398	rent	food	course fee
12863	health	food	lunch
718	rent	food	test
4561	education	food	unknown
26044	entertainment	food	kolkatagrocery shopping
27815	entertainment	food	gift
8354	utilities	food	monthly rent payment

```
In [59]: #Hyperparameter tuning for Naive Bayes + TF-IDF

text_col = "notes"
cat_cols = ["payment_mode", "location", "transaction_type"]
num_cols = ["amount"]

def extract_date_features(df_in):
    df2 = df_in.copy()
    dates = pd.to_datetime(df2["date"], errors="coerce")
    df2["dt_year"] = dates.dt.year.fillna(0).astype(int)
    df2["dt_month"] = dates.dt.month.fillna(0).astype(int)
    df2["dt_day"] = dates.dt.day.fillna(0).astype(int)
    df2["dt_dow"] = dates.dt.dayofweek.fillna(0).astype(int)
    return df2[["dt_year", "dt_month", "dt_day", "dt_dow"]]

date_transformer = Pipeline([
    ("extract", FunctionTransformer(extract_date_features, validate=False)),
    ("impute", SimpleImputer(strategy="constant", fill_value=0)),
    ("scale", MinMaxScaler())
])

text_transformer = TfidfVectorizer(stop_words="english")

cat_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="constant", fill_value="unknown")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

num_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="median")),
    ("scale", MinMaxScaler())
])

preprocessor = ColumnTransformer(
    transformers=[
        ("text", text_transformer, text_col),
        ("cat", cat_transformer, cat_cols),
        ("num", num_transformer, num_cols),
        ("date", date_transformer, ["date"])
    ],
    remainder="drop"
)

pipeline = Pipeline([
    ("pre", preprocessor),
    ("clf", MultinomialNB())
])

#Hyperparameters to search
params = {
    "pre_text_max_features": [5000, 10000, 20000],
    "pre_text_ngram_range": [(1,1), (1,2)],
    "clf_alpha": [0.1, 0.5, 1.0]
}

grid = GridSearchCV(
    pipeline,
    params,
    cv=3,
    scoring="accuracy",
    n_jobs=-1,
    verbose=1
)

grid.fit(X_train, y_train)

print("Best Params:", grid.best_params_)
print("Best CV Accuracy:", grid.best_score_)

#Final model
best_model = grid.best_estimator_

#Evaluate on test set
y_pred = best_model.predict(X_test)
print("Test Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred, zero_division=0))

#Save final model
joblib.dump(best_model, "output/nb_tuned_pipeline.joblib")
print("Saved: output/nb_tuned_pipeline.joblib")
```

```
Fitting 3 folds for each of 18 candidates, totalling 54 fits
Best Params: {'clf_alpha': 1.0, 'pre_text_max_features': 5000, 'pre_text_ngram_range': (1, 1)}
Best CV Accuracy: 0.2675129111950188
Test Accuracy: 0.26369235891459963

      precision    recall   f1-score   support

        bonus      0.25     0.17     0.20      234
      education      0.00     0.00     0.00      266
  entertainment      0.00     0.00     0.00      472
         food      0.27     0.91     0.42     1391
      freelance      0.21     0.10     0.13      93
       health      0.00     0.00     0.00     164
  investment      0.15     0.07     0.09      91
       misc      0.00     0.00     0.00      9
other income      0.32     0.65     0.43     149
      others      0.20     0.17     0.18     254
       rent      0.24     0.08     0.12     1131
      salary      0.23     0.14     0.17     232
     savings      0.00     0.00     0.00     195
      travel      0.04     0.00     0.00     676
     unknown      0.00     0.00     0.00      86
 utilities      0.00     0.00     0.00     564

  accuracy          0.26      --      6007
macro avg      0.12     0.14     0.11      6007
weighted avg      0.15     0.26     0.15      6007
```

Saved: output/nb_tuned_pipeline.joblib

```
In [60]: #Train Logistic Regression baseline (strong model)

text_col = "notes"
cat_cols = ["payment_mode", "location", "transaction_type"]
num_cols = ["amount"]

#date feature extractor
def extract_date_features(df_in):
    df2 = df_in.copy()
    dates = pd.to_datetime(df2["date"], errors="coerce")
    df2["dt_year"] = dates.dt.year.fillna(0).astype(int)
    df2["dt_month"] = dates.dt.month.fillna(0).astype(int)
    df2["dt_day"] = dates.dt.day.fillna(0).astype(int)
    df2["dt_dow"] = dates.dt.dayofweek.fillna(0).astype(int)
    return df2[["dt_year", "dt_month", "dt_day", "dt_dow"]]

date_transformer = Pipeline([
    ("extract", FunctionTransformer(extract_date_features, validate=False)),
    ("impute", SimpleImputer(strategy="constant", fill_value=0)),
    ("scale", MinMaxScaler())
])

#text + categorical + numeric transformers
text_transformer = TfidfVectorizer(max_features=20000, stop_words="english")

cat_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="constant", fill_value="unknown")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

num_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="median")),
    ("scale", MinMaxScaler())
])

#full preprocessor
preprocessor = ColumnTransformer(
    transformers=[
        ("text", text_transformer, text_col),
        ("cat", cat_transformer, cat_cols),
        ("num", num_transformer, num_cols),
        ("date", date_transformer, ["date"])
    ],
    remainder="drop"
)

#Logistic Regression model
lr = Pipeline([
    ("pre", preprocessor),
    ("clf", LogisticRegression(
        max_iter=2000,
        n_jobs=-1,
        class_weight="balanced",
        solver="lbfgs",
        multi_class="auto"
    ))
])

#Train
lr.fit(X_train, y_train)

#Evaluate
y_pred_lr = lr.predict(X_test)
print("LR Accuracy:", accuracy_score(y_test, y_pred_lr))
print(classification_report(y_test, y_pred_lr, zero_division=0))

#Save
joblib.dump(lr, "output/lr_pipeline.joblib")
print("Saved: output/lr_pipeline.joblib")
```

```
LR Accuracy: 0.10071583152988181
      precision    recall   f1-score   support
      bonus        0.35     0.10     0.16     234
      education     0.08     0.17     0.11     266
      entertainment  0.13     0.05     0.07     472
      food          0.31     0.08     0.13    1391
      freelance     0.19     0.41     0.26      93
      health         0.04     0.22     0.07    164
      investment     0.19     0.44     0.27     91
      misc           0.00     0.22     0.00      9
other income     0.31     0.72     0.43    149
      others         0.08     0.05     0.06    254
      rent           0.19     0.05     0.08   1131
      salary          0.27     0.09     0.14    232
      savings         0.05     0.27     0.09    195
      travel          0.11     0.01     0.03    676
      unknown         0.01     0.07     0.02     86
      utilities       0.11     0.04     0.06    564
accuracy          -        -       0.10    6007
macro avg        0.15     0.19     0.12    6007
weighted avg      0.19     0.10     0.10    6007
```

Saved: output/lr_pipeline.joblib

In [61]: #Logistic Regression WITHOUT notes (structured features only)

```
#Columns to use
cat_cols = ["payment_mode", "location", "transaction_type"]
num_cols = ["amount"]

#date feature extractor
def extract_date_features(df_in):
    df2 = df_in.copy()
    dates = pd.to_datetime(df2["date"], errors="coerce")
    df2["dt_year"] = dates.dt.year.fillna(0).astype(int)
    df2["dt_month"] = dates.dt.month.fillna(0).astype(int)
    df2["dt_day"] = dates.dt.day.fillna(0).astype(int)
    df2["dt_dow"] = dates.dt.dayofweek.fillna(0).astype(int)
    return df2[["dt_year", "dt_month", "dt_day", "dt_dow"]]

date_transformer = Pipeline([
    ("extract", FunctionTransformer(extract_date_features, validate=False)),
    ("impute", SimpleImputer(strategy="constant", fill_value=0)),
    ("scale", MinMaxScaler())
])

#categorical + numeric transformers
cat_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="constant", fill_value="unknown")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

num_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="median")),
    ("scale", MinMaxScaler())
])

#full preprocessor (No text transformer)
preprocessor = ColumnTransformer(
    transformers=[
        ("cat", cat_transformer, cat_cols),
        ("num", num_transformer, num_cols),
        ("date", date_transformer, ["date"]),
    ],
    remainder="drop"
)

#Logistic Regression
lr_structured = Pipeline([
    ("pre", preprocessor),
    ("clf", LogisticRegression(
        max_iter=2000,
        class_weight="balanced",
        n_jobs=-1
    ))
])

#Train
lr_structured.fit(X_train, y_train)

#Evaluate
y_pred_lr_struct = lr_structured.predict(X_test)
print("Structured LR Accuracy:", accuracy_score(y_test, y_pred_lr_struct))
print(classification_report(y_test, y_pred_lr_struct, zero_division=0))

#Save
joblib.dump(lr_structured, "output/lr_structured_pipeline.joblib")
print("Saved: output/lr_structured_pipeline.joblib")
```

```
Structured LR Accuracy: 0.08123855501914433
precision    recall   f1-score   support
bonus        0.00     0.00     0.00      234
education    0.05     0.08     0.06      266
entertainment 0.12     0.04     0.06      472
food         0.29     0.01     0.02     1391
freelance    0.20     0.52     0.29      93
health        0.04     0.27     0.07     164
investment    0.20     0.37     0.26      91
misc          0.00     0.67     0.01       9
other income  0.31     0.97     0.47     149
others        0.00     0.00     0.00     254
rent          0.21     0.07     0.11    1131
salary        0.24     0.02     0.03     232
savings       0.05     0.24     0.08     195
travel         0.10     0.01     0.03     676
unknown       0.01     0.01     0.01      86
utilities     0.13     0.02     0.04     564
accuracy      0.08
macro avg     0.12     0.21     0.10     6007
weighted avg  0.17     0.08     0.06     6007
```

Saved: output/lr_structured_pipeline.joblib

In [62]: #Train Random Forest Classifier (handles noisy data better)

```
#Feature groups
text_col = "notes"
cat_cols = ["payment_mode", "location", "transaction_type"]
num_cols = ["amount"]

#Date features
def extract_date_features(df_in):
    df2 = df_in.copy()
    dates = pd.to_datetime(df2["date"], errors="coerce")
    df2["dt_year"] = dates.dt.year.fillna(0).astype(int)
    df2["dt_month"] = dates.dt.month.fillna(0).astype(int)
    df2["dt_day"] = dates.dt.day.fillna(0).astype(int)
    df2["dt_dow"] = dates.dt.dayofweek.fillna(0).astype(int)
    return df2[["dt_year", "dt_month", "dt_day", "dt_dow"]]

date_transformer = Pipeline([
    ("extract", FunctionTransformer(extract_date_features, validate=False)),
    ("impute", SimpleImputer(strategy="constant", fill_value=0)),
    ("scale", MinMaxScaler())
])

#Text transformer
text_transformer = Pipeline([
    ("tfidf", TfidfVectorizer(max_features=20000, stop_words="english"))
])

#Categorical + Numeric transformers
cat_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="constant", fill_value="unknown")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

num_transformer = Pipeline([
    ("impute", SimpleImputer(strategy="median")),
    ("scale", MinMaxScaler())
])

#Full preprocessor
preprocessor = ColumnTransformer(
    transformers=[
        ("text", text_transformer, text_col),
        ("cat", cat_transformer, cat_cols),
        ("num", num_transformer, num_cols),
        ("date", date_transformer, ["date"]),
    ],
    remainder="drop"
)

#Random Forest model
rf = Pipeline([
    ("pre", preprocessor),
    ("clf", RandomForestClassifier(
        n_estimators=300,
        random_state=42,
        class_weight="balanced",
        n_jobs=-1
    ))
])

#Train
rf.fit(X_train, y_train)

#Evaluate
y_pred_rf = rf.predict(X_test)
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf, zero_division=0))

#Save
joblib.dump(rf, "output/rf_pipeline.joblib")
print("Saved: output/rf_pipeline.joblib")
```

```
Random Forest Accuracy: 0.20559347428000666
      precision    recall   f1-score   support
      bonus        0.25     0.24     0.24      234
      education    0.07     0.06     0.06      266
      entertainment 0.10     0.08     0.09      472
      food         0.26     0.32     0.29     1391
      freelance    0.22     0.24     0.23      93
      health        0.01     0.01     0.01     164
      investment    0.18     0.16     0.17      91
      misc          0.00     0.00     0.00       9
other income    0.35     0.54     0.42     149
      others        0.12     0.13     0.12     254
      rent          0.32     0.30     0.31    1131
      salary        0.26     0.21     0.23     232
      savings       0.03     0.05     0.04     195
      travel         0.15     0.12     0.14     676
      unknown        0.01     0.01     0.01      86
      utilities     0.11     0.09     0.10     564
      accuracy      0.21      --      6007
      macro avg     0.15     0.16     0.15     6007
      weighted avg   0.20     0.21     0.20     6007
```

Saved: output/rf_pipeline.joblib

```
In [63]: import nbformat
from nbconvert.preprocessors import ExecutePreprocessor

def run_notebook(path):
    with open(path, "r", encoding="utf-8") as f:
        nb = nbformat.read(f, as_version=4)

    ep = ExecutePreprocessor(timeout=None, kernel_name="python3")
    ep.preprocess(nb)

    print(f"{path} ran successfully")

run_notebook("predict_code.ipynb")
run_notebook("app_code.ipynb")
```

predict_code.ipynb ran successfully
app_code.ipynb ran successfully